

This project focuses on building various machine learning models for Heart Disease prediction

```
In [1]: ## Importing the necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import janitor
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.preprocessing import LabelEncoder
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: ## Load the dataset
df = pd.read_csv("C:/Users/ADMIN/Desktop/Data Science/Datasets/Datasets/heart.csv")
```

```
In [3]: ## View the first few observations of the dataset
df.head(5)
```

```
Out[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	tb
--	-----	-----	----	----------	------	-----	---------	---------	-------	---------	-------	----	------	----

0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	



```
In [4]: ## Assess the structure of the dataset
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1025 non-null   int64
 1   sex         1025 non-null   int64
 2   cp          1025 non-null   int64
 3   trestbps    1025 non-null   int64
 4   chol        1025 non-null   int64
 5   fbs         1025 non-null   int64
 6   restecg     1025 non-null   int64
 7   thalach     1025 non-null   int64
 8   exang       1025 non-null   int64
 9   oldpeak     1025 non-null   float64
10   slope       1025 non-null   int64
11   ca          1025 non-null   int64
12   thal        1025 non-null   int64
13   target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB

```

```

In [5]: ## Checking for missing values
df.isna().sum()

```

```

Out[5]: age         0
sex         0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64

```

```

In [6]: ## Check for duplicates
df.duplicated().sum()

```

```

Out[6]: 723

```

```

In [7]: ## Number of males and females whose heart data is stored in the dataset
df.sex.value_counts()

```

```

Out[7]: sex
1      713
0      312
Name: count, dtype: int64

```


```

In [8]: ## Count of the number of males and females who have heart disease
df.sex[df.target==1].value_counts()

```

```
Out[8]: sex
1      300
0      226
Name: count, dtype: int64
```

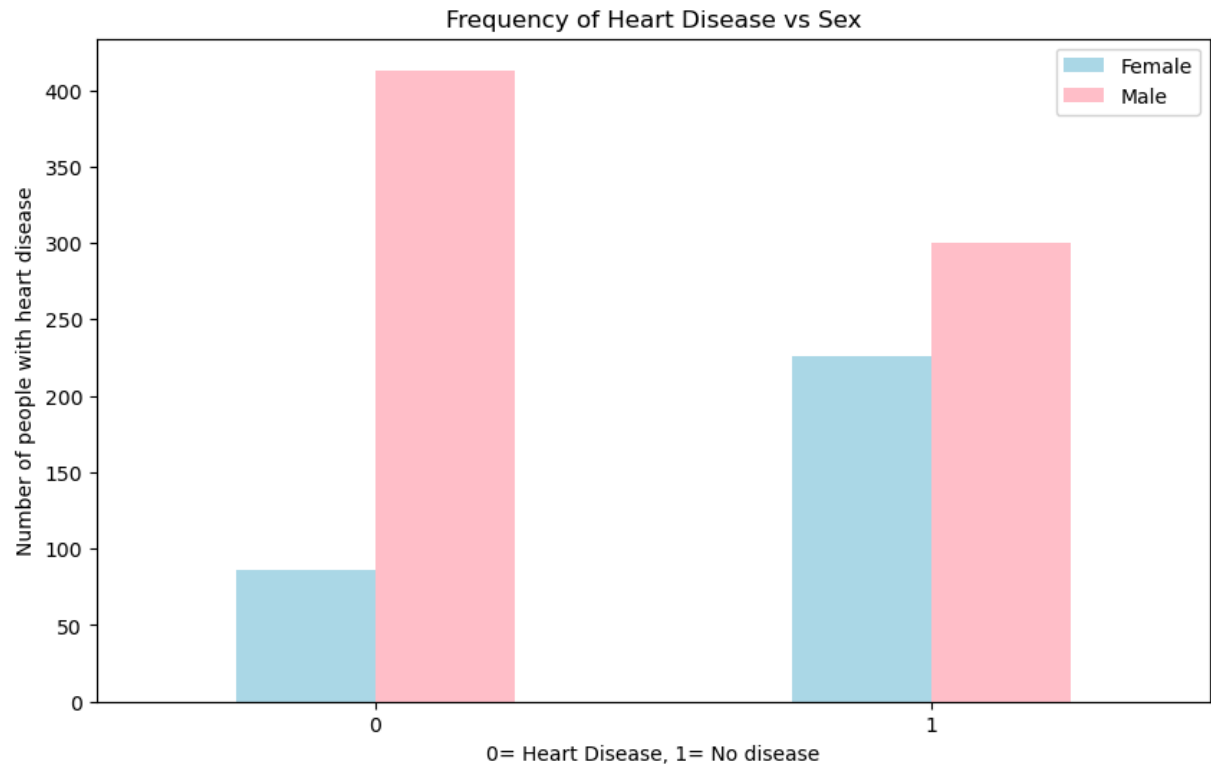
```
In [9]: ## Visualize your results
df.sex[df.target==1].value_counts().plot(kind='bar',figsize=(10,6),color=['green','blue'])
plt.title("Count of the number of males and females with heart disease")
plt.xticks([0,1], labels = ['Female', 'Male'], rotation =0)
plt.show()
```



```
In [10]: ## Contingency table
pd.crosstab(df.target,df.sex)
```

```
Out[10]:    sex    0    1
target
0      86  413
1      226  300
```

```
In [11]: ## Visualize the contingency table
pd.crosstab(df.target,df.sex).plot(kind='bar',figsize=(10,6),color=["lightblue","lightgreen"])
plt.title("Frequency of Heart Disease vs Sex")
plt.xlabel("0= Heart Disease, 1= No disease")
plt.ylabel("Number of people with heart disease")
plt.legend(["Female", "Male"])
plt.xticks(rotation=0);
```




```
In [12]: ## Building a Correlation Matrix
df.corr()
```

Out[12]:

	age	sex	cp	trestbps	chol	fbs	restecg	thal
age	1.000000	-0.103240	-0.071966	0.271121	0.219823	0.121243	-0.132696	-0.390227
sex	-0.103240	1.000000	-0.041119	-0.078974	-0.198258	0.027200	-0.055117	-0.049365
cp	-0.071966	-0.041119	1.000000	0.038177	-0.081641	0.079294	0.043581	0.306839
trestbps	0.271121	-0.078974	0.038177	1.000000	0.127977	0.181767	-0.123794	-0.039264
chol	0.219823	-0.198258	-0.081641	0.127977	1.000000	0.026917	-0.147410	-0.021772
fbs	0.121243	0.027200	0.079294	0.181767	0.026917	1.000000	-0.104051	-0.008866
restecg	-0.132696	-0.055117	0.043581	-0.123794	-0.147410	-0.104051	1.000000	0.048411
thalach	-0.390227	-0.049365	0.306839	-0.039264	-0.021772	-0.008866	0.048411	1.000000
exang	0.088163	0.139157	-0.401513	0.061197	0.067382	0.049261	-0.065606	-0.380137
oldpeak	0.208137	0.084687	-0.174733	0.187434	0.064880	0.010859	-0.050114	-0.349752
slope	-0.169105	-0.026666	0.131633	-0.120445	-0.014248	-0.061902	0.086086	0.395195
ca	0.271551	0.111729	-0.176206	0.104554	0.074259	0.137156	-0.078072	-0.207195
thal	0.072297	0.198424	-0.163341	0.059276	0.100244	-0.042177	-0.020504	-0.098001
target	-0.229324	-0.279501	0.434854	-0.138772	-0.099966	-0.041164	0.134468	0.422574

```
In [13]: ## Visualize the correlation matrix
cor_mat=df.corr()
plt.figure(figsize=(10,10))
sns.heatmap(cor_mat,annot=True,cmap="coolwarm", fmt = ".2f")
plt.title("Correlation Matrix")
plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.show()
```



Machine Learning

```
In [14]: ## Standardization using the MinMax Scaling
from sklearn.preprocessing import MinMaxScaler
scal=MinMaxScaler()
feat=['age', 'trestbps', 'thalach', 'oldpeak', 'chol']
```

```
df[feat] = scal.fit_transform(df[feat])
df.head()
```

Out[14]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope
0	0.479167	1	0	0.292453	0.196347	0	1	0.740458	0	0.161290	2
1	0.500000	1	0	0.433962	0.175799	1	0	0.641221	1	0.500000	0
2	0.854167	1	0	0.481132	0.109589	0	1	0.412214	1	0.419355	0
3	0.666667	1	0	0.509434	0.175799	0	1	0.687023	0	0.000000	2
4	0.687500	0	0	0.415094	0.383562	1	1	0.267176	0	0.306452	1

In [15]: *## Creating Features and Target variable*
`X=df.drop("target",axis=1).values`
`Y=df.target.values`

In [16]: *## Splitting the data into train and test sets*
`from sklearn.model_selection import train_test_split`
`X_train,X_test,Y_train,Y_test=train_test_split(X,Y,random_state=42,test_size=0.2)`

In [17]: *## Create a function for evaluating metrics*
`from sklearn.metrics import accuracy_score,recall_score,f1_score,precision_score,ro`
`def evaluation(Y_test,Y_pred):`
 `acc=accuracy_score(Y_test,Y_pred)`
 `rcl=recall_score(Y_test,Y_pred)`
 `f1=f1_score(Y_test,Y_pred)`

 `metric_dict={'accuracy': round(acc,3),`
 `'recall': round(rcl,3),`
 `'F1 score': round(f1,3),`

 `}`

 `return print(metric_dict)`

Fitting and Comparing different Models

In [18]: *## K-Nearest Neighbors*
`np.random.seed(42)`
`from sklearn.neighbors import KNeighborsClassifier`
`Knn_clf= KNeighborsClassifier()`
`Knn_clf.fit(X_train,Y_train)`
`Knn_Y_pred=Knn_clf.predict(X_test)`
`Knn_score=Knn_clf.score(X_test,Y_test)`
`#print(Knn_score)`
`evaluation(Y_test,Knn_Y_pred)`

{'accuracy': 0.859, 'recall': 0.883, 'F1 score': 0.863}

```
In [19]: ## Logistic Regression
np.random.seed(42)
from sklearn.linear_model import LogisticRegression
LR_clf=LogisticRegression()
LR_clf.fit(X_train,Y_train)
LR_Y_pred=LR_clf.predict(X_test)
LR_score=LR_clf.score(X_test,Y_test)
#print(LR_score)
evaluation(Y_test,LR_Y_pred)
```

```
{'accuracy': 0.8, 'recall': 0.874, 'F1 score': 0.814}
```

```
In [20]: ## Random Forest
np.random.seed(42)
from sklearn.ensemble import RandomForestClassifier
RF_clf=RandomForestClassifier(n_estimators=450)
RF_clf.fit(X_train,Y_train)
RF_score=RF_clf.score(X_test,Y_test)
RF_Y_pred=RF_clf.predict(X_test)
#print(RF_score)
evaluation(Y_test,RF_Y_pred)
```

```
{'accuracy': 0.985, 'recall': 0.971, 'F1 score': 0.985}
```

```
In [21]: ## Support Vector Machines
np.random.seed(42)
from sklearn.svm import SVC
SVC_clf=SVC()
SVC_clf.fit(X_train,Y_train)
SVC_score=SVC_clf.score(X_test,Y_test)
SVC_Y_pred=SVC_clf.predict(X_test)
#print(SVC_score)
evaluation(Y_test,SVC_Y_pred)
```

```
{'accuracy': 0.883, 'recall': 0.932, 'F1 score': 0.889}
```

Comparing the performance of different models

```
In [22]: ## Model comparison
model_comp = pd.DataFrame({'Model': ['Logistic Regression','Random Forest',
                                     'K-Nearest Neighbour','Support Vector Machine'], 'Accuracy': [LR_score*100,Knn_score*100,SVC_score*100]})
model_comp
```

```
Out[22]:
```

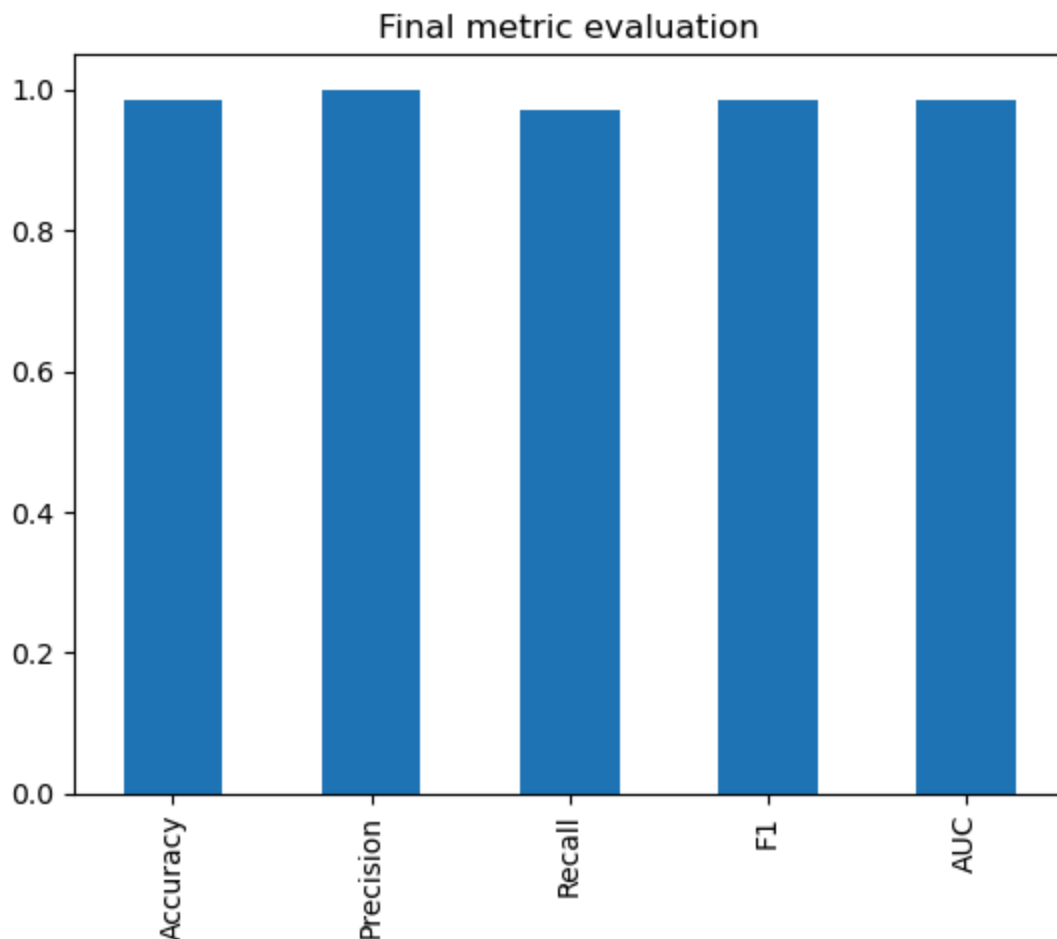
	Model	Accuracy
0	Logistic Regression	80.000000
1	Random Forest	98.536585
2	K-Nearest Neighbour	85.853659
3	Support Vector Machine	88.292683

```
In [26]: print(" Best evaluation parameters achieved with Random Forest:")  
evaluation(Y_test, RF_Y_pred)
```

Best evaluation parameters achieved with Random Forest:
{'accuracy': 0.985, 'recall': 0.971, 'F1 score': 0.985}

Visualization of the Models Results


```
In [27]: final_metrics={'Accuracy': RF_clf.score(X_test,Y_test),  
                        'Precision': precision_score(Y_test,RF_Y_pred),  
                        'Recall': recall_score(Y_test,RF_Y_pred),  
                        'F1': f1_score(Y_test,RF_Y_pred),  
                        'AUC': roc_auc_score(Y_test,RF_Y_pred)}  
  
metrics=pd.DataFrame(final_metrics,index=[0])  
  
metrics.T.plot.bar(title='Final metric evaluation',legend=False);
```



Create the confusion matrix of the best model


```
In [28]: from sklearn.metrics import confusion_matrix

fig,ax=plt.subplots()
ax=sns.heatmap(confusion_matrix(Y_test,RF_Y_pred),annot=True,cbar=True);
```



Let's save our model using pickle

```
In [30]: ## Lets save our model using pickle
import pickle as pkl
pkl.dump(RF_clf,open("finali_model.p","wb"))
```